

PETERSON'S SOLUTION

Peterson's solution is a classic **Software-Based Solution** to the critical-section problem.

Peterson's solution is restricted to two processes that alternate execution between their critical sections and remainder sections.

The processes are numbered ***P0*** and ***P1***. Let ***Pi*** represents one process and ***Pj*** represents other processes (i.e. $j = i-1$)

do

{

```
flag[i] = true;
turn = j;
while (flag[j] && turn == j);
```

Critical Section

```
flag[i] = false;
```

Remainder Section

} while (true);

Peterson's solution requires the two processes to share two data items:

int turn;

boolean flag[2];

The variable turn indicates whose turn it is to enter its critical section. At any point of time the turn value will be either 0 or 1 but not both.

- if $turn == i$, then process ***Pi*** is allowed to execute in its critical section.
- if $turn == j$, then process ***Pj*** is allowed to execute in its critical section.
- The flag array is used to indicate if a process is ready to enter its critical section.

Example: if $flag[i]$ is true, this value indicates that ***Pi*** is ready to enter its critical section.

- To enter the critical section, process ***Pi*** first sets **$flag[i]=true$** and then sets **$turn=j$** , thereby ***Pi*** checks if the other process wishes to enter the critical section, it can do so.
- If both processes try to enter at the same time, turn will be set to both i and j at the same time. Only one of these assignments will be taken. The other will occur but will be overwritten immediately.
- The eventual value of turn determines which of the two processes is allowed to enter its critical section first.

The above code must satisfy the following requirements:

1. Mutual exclusion
2. The progress
3. The bounded-waiting

Check for Mutual Exclusion

- Each ***Pi*** enters its critical section only if either **$flag[j] == false$** or **$turn == i$** .
- If both processes can be executing in their critical sections at the same time, then **$flag[0] == flag[1] == true$** . But the value of turn can be either **0** or **1** but cannot be both.

- Hence P_0 and P_1 could not have successfully executed their **while** statements at about the same time.
- If P_i executed "**turn == j**" and the process P_j executed **flag[j]=true** then P_j will have successfully executed the while statement. Now P_j will enter into its **Critical section**.
- At this time, **flag[j] == true** and **turn == j** and this condition will persist as long as P_j is in its critical section. As a result, mutual exclusion is preserved.

Check for Progress and Bounded-waiting

- The while loop is the only possible way that a process P_i can be prevented from entering the critical section only if it is stuck in the while loop with the condition **flag[j] == true** and **turn == j**.
- If P_j is not ready to enter the critical section, then **flag[j] == false** and P_i can enter its critical section.
- If P_j has set **flag[j] == true** and is also executing in its while statement, then either **turn == i** or **turn == j**.
- If **turn == i**, then P_i will enter the critical section. If **turn == j**, then P_j will enter the critical section.
- Once P_j exits its critical section, it will reset **flag[j]** to false, allowing P_i to enter its critical section.
- If P_j resets **flag[j]** to true, it must also set **turn == i**. Thus, since P_i does not change the value of the variable **turn** while executing the while statement, P_i will enter the critical section (**Progress**) after at most one entry by P_j (**Bounded Waiting**).

Problem with Peterson Solution

There are no guarantees that Peterson's solution will work correctly on modern computer architectures perform basic machine-language instructions such as load and store.

- The critical-section problem could be solved simply in a single-processor environment if we could prevent interrupts from occurring while a shared variable was being modified. This is a Non-preemptive kernel approach.
- We could be sure that the current sequence of instructions would be allowed to execute in order without preemption.
- No other instructions would be run, so no unexpected modifications could be made to the shared variable.

This solution is not as feasible in a multiprocessor environment.

- Disabling interrupts on a multiprocessor can be time consuming, since the message is passed to all the processors.
- This message passing delays entry into each critical section and system efficiency decreases and if the clock is kept updated by interrupts there will be an effect on a system's clock.